

Processes

Tran Giang Son, tran-giang.son@usth.edu.vn

ICT Department, USTH

Contents

- Process
- Process Scheduling
- Process Creation

3 / 63

Program \neq Process

Program \neq Process

- Program is **passive**,

Program \neq Process

- Program is **passive**, stored on disk as an **executable file**
 - e.g. /bin/ls


```
$ ls -la /bin/ls
-rwxr-xr-x 1 root root 118280 Mar 14 2015 /bin/ls
```


- Process is a program in execution state (**active**)
- Created by the system or a parent process
- Uniquely identified (PID)

Process States

```
graph LR; Start(( )) --> new((new)); new -- admitted --> ready((ready)); ready -- scheduled --> running((running)); running -- interrupt --> ready; running -- "event wait" --> waiting((waiting)); waiting -- "event finish" --> ready; running -- exit --> terminated((terminated)); terminated --> End((( )))
```

Process States



- **new**: process has just been created
- **ready**: waiting to be assigned (scheduled) to a processor
- **running**: it's executing instructions
- **waiting**: waiting for some events to occur
- **terminated**: finished execution

Process Control Block

- Represents a process
- Stored in memory
- Not accessible to process, only for kernel's process schedulers (later)


```

/*
 * If the new process paused because it was
 * swapped out, set the stack level to the last call
 * to savu(u_ssav). This means that the return
 * which is executed immediately after the call to aretu
 * actually returns from the last routine which did
 * the savu.
 *
 * You are not expected to understand this.
 */
if(rp->p_flag&SSWAP) {
    rp->p_flag =& ~SSWAP;
    aretu(u.u_ssav);
}

```

Source: [Unix v6 Source Code](#), line 2230-2243

Process Scheduling

- Multiple processes running at the same time.
- Process scheduler is a part that decides which processes to be executed at a certain time.

Process Scheduling Types

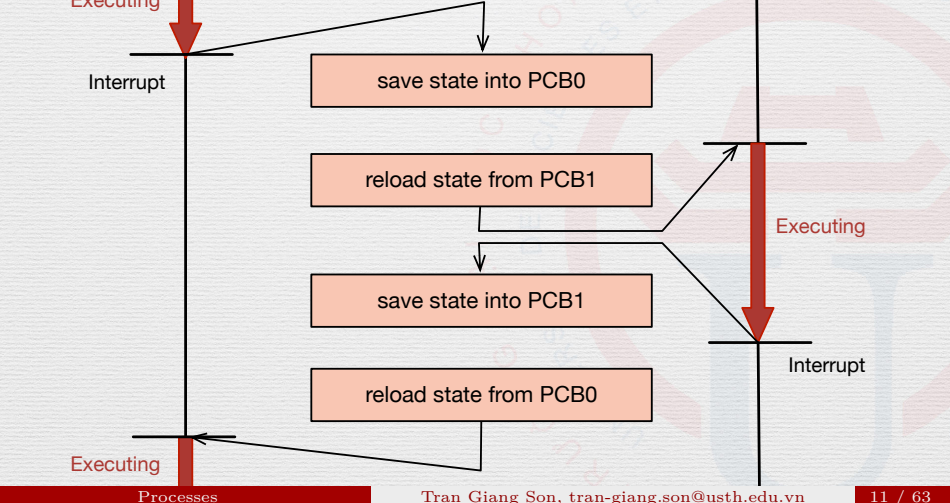
- Pause running processes
 - Preemption: OS forcibly pauses running processes
 - Non-preemption (also cooperation): processes willing to pause itself
- Duration between each «switch»
 - Short term scheduler: milliseconds (fast, responsive)
 - Long term scheduler: seconds/minutes (batch jobs)

Context Switch

Process A Operating System Process B

Executing

```
graph LR; A[Process A] -- "Executing" --> OS[Operating System]; OS --> B[Process B];
```



Context Switch

- Switch between processes
 - Save data of old process
 - Load previously saved data of new process
- Context switch is overhead
 - No work done for processes during context switch
 - Time slice (time between each switch) is hardware-limited

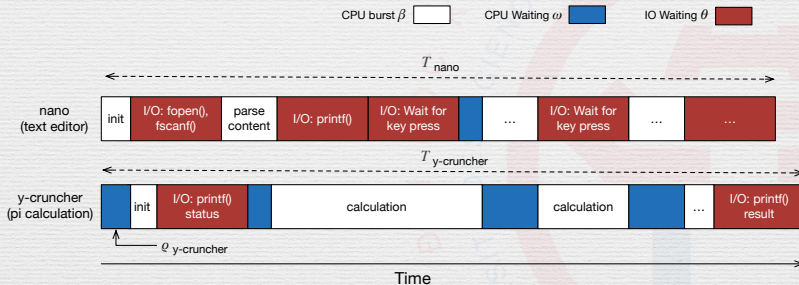
Scheduling Algorithms

Scheduling Concepts

- Timing
 - **Burst time** β : CPU time needed for a short task
 - **CPU Waiting time** ω : time a process in the **ready** queue
 - **I/O time** θ : time a process in the **waiting** queue
 - **Turnaround time** T : total time to execute a process, from start to end, including waiting times (in **ready** and **waiting** queues)
 - **Response time** ρ : time from a submitted request until first response

Scheduling Concepts

**Burst time β , CPU Waiting time ω , I/O time θ ,
Turnaround time T , Response time ρ .**



$$T = \sum \beta_i + \sum \omega_j + \sum \theta_k$$

Scheduling Concepts

- Others
 - **CPU utilization:** percentage of CPU usage
 - **Throughput:** number of processes completing their execution per time unit
 - **Priority:** an integer number for each process, indicating its importance

Scheduling Goals

- Idle = waste of energy, so...
- Maximize
 - CPU utilization
 - Throughput
- Minimize
 - Turnaround time: for calculation
 - **Waiting time**: for typical desktop systems
 - Response time: for server systems


Algorithms

- **First Come, First Served**
- Shortest-Job-First
- Shortest-Remaining-Time-First
- Round Robin
- Multilevel Queue
- Multilevel Feedback Queue


First Come, First Served: What?

- **Non-preemptive**
- **Non-priority**
- Allocate CPU to next process in ready queue
 - **based on the process arrival time**
- Wait till it finishes CPU usage
 - I/O wait
 - Voluntarily stop
- Switch to next process

First Come, First Served: How?



First Come, First Served: How?




$$\omega_{firefox} = 0$$
$$\omega_{gedit} = 23$$
$$\omega_{thunderbird} = 24$$

$$\bar{\omega}_{fcfs} = \frac{0+23+24}{3} = 15.67$$

First Come, First Served: How?


If arrival time is reversed?

Arrival time (ms)	Process	Burst time (ms)
0	thunderbird	4
1	gedit	3
3	firefox	24



First Come, First Served: How?

If arrival time is reversed?



$$\omega_{thunderbird} = 0$$

$$\omega_{gedit} = 3$$

$$\omega_{firefox} = 4$$

$$\bar{\omega}_{fcfs} = \frac{0+3+4}{3} = 2.33$$

Algorithms


- First Come, First Served
- **Shortest-Job-First**
- Shortest-Remaining-Time-First
- Round Robin
- Multilevel Queue
- Multilevel Feedback Queue

Shortest-Job-First: What?


- **Non-preemptive**
- **Priority**
- Allocate CPU to next process in ready queue
 - NOT based on the process arrival time
 - based on the estimated CPU burst
 - **choose the process with shortest estimated burst**
 - If there are two processes with shortest bursts, FCFS
 - That's priority
- Optimal for waiting time \bar{w}_{sjf}
- Problem: processes with long CPU bursts will rarely be scheduled

Shortest-Job-First: How?

Arrival time (ms)	Process	Burst time (ms)
0	firefox	24
0	gedit	3
0	thunderbird	4



Shortest-Job-First: How?



$$\omega_{gedit} = 0$$

$$\omega_{thunderbird} = 3$$

$$\omega_{firefox} = 7$$

$$\overline{\omega}_{sjf} = \frac{0+3+7}{3} = 3.33$$

Algorithms

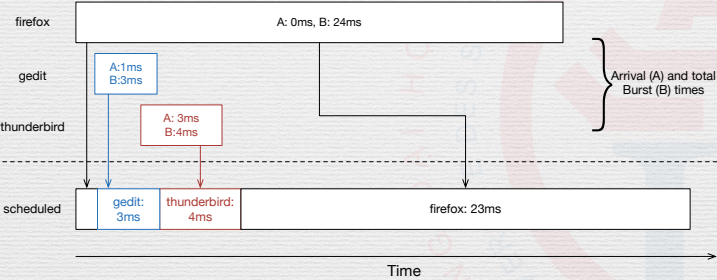
- First Come, First Served
- Shortest-Job-First
- **Shortest-Remaining-Time-First**
- Round Robin
- Multilevel Queue
- Multilevel Feedback Queue

Shortest-Remaining-Time-First: What?

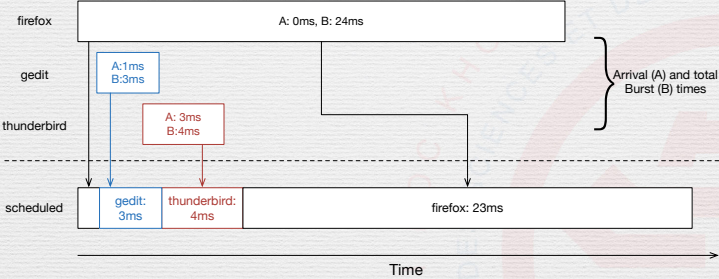
- **Preemptive**
- **Priority**
- Each preemption, allocate CPU to next process in ready queue
 - NOT based on the process arrival time
 - based on the estimate **remaining** CPU burst
 - choose the process with shortest **remaining** burst
 - If there are two processes with shortest remaining bursts, FCFS
 - That's priority
- Also called preemptive Shortest-Job-First

Shortest-Remaining-Time-First: Example

Arrival time (ms)	Process	Burst time (ms)
0	firefox	24
1	gedit	3
3	thunderbird	4



Shortest-Remaining-Time-First: Example



$$\omega_{firefox} = 0 + 7$$

$$\omega_{gedit} = 0$$

$$\omega_{thunderbird} = 1$$

$$\overline{\omega}_{srtf} = \frac{7+0+1}{3} = 2.33$$

Algorithms

- First Come, First Served
- Shortest-Job-First
- Shortest-Remaining-Time-First
- **Round Robin**
- Multilevel Queue
- Multilevel Feedback Queue


Round Robin: What?

- **Preemptive**
- **Non-priority**
- Similar to FCFS, but add periodical preemption
 - Each process has a time slice (duration)
 - Fixed, or
 - Dynamic
 - After time slice finishes
 - Process is preempted
 - Put to end of ready queue

Round Robin: What?


- Good for interactive systems
- Better average response time ρ
- n : number of processes in ready queue
- s : time slice (also called quantum)
- \rightarrow Waiting time **per round**

$$\omega_r = (n - 1) * s$$

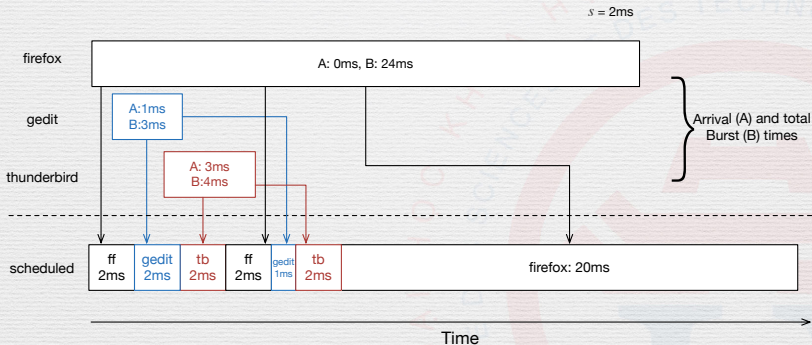


Round Robin: Example

Arrival time (ms)	Process	Burst time (ms)
0	firefox	24
1	gedit	3
3	thunderbird	4



Round Robin: Example



$$\omega_{firefox} = 0 + 4 + 3 = 7$$

$$\omega_{gedit} = 1 + 4 = 5$$

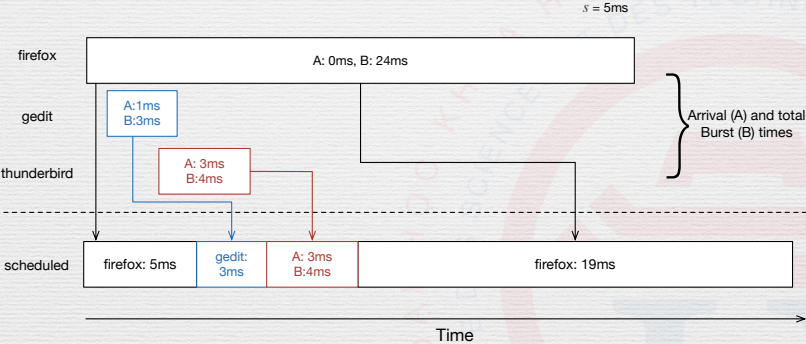
$$\omega_{thunderbird} = 1 + 3 = 4$$

$$\bar{\omega}_{rr} = \frac{7+5+4}{3} = 5.33$$

Round Robin: Time Slice (s)

- Smaller s , more context switches
- Too small s , too many preemption (\sim context switches)
 - Remind: context switch does nothing productive
 - Overhead
- Too large s , close to FCFS

Round Robin: Time Slice (s)



- Only 3 context switches total
- 2 context switches for firefox

Algorithms

- First Come, First Served
- Shortest-Job-First
- Shortest-Remaining-Time-First
- Round Robin
- **Multilevel Queue**
- Multilevel Feedback Queue

Multilevel Queue: What?

- Combination of previous algorithms
- Ready queue is split into several subqueues
 - Based on characteristics of processes
 - **Permanently** assign them to different subqueues
 - Example: building Linux kernel while surfing web and listening to music
- Each queue uses different scheduling algorithm
 - Example: RR for foreground subqueue, FCFS for background subqueue.

Multilevel Queue: What?

- Problem: different subqueues, choose which subqueue to schedule?
 - System processes: sshd, postfix, ...
 - Interactive processes: firefox, chrome, ...
 - Background processes: gcc, make, wget, ...
- Solution
 - Subqueue priority
 - Subqueue time slice

Multilevel Queue: Subqueue Priority

- Category each subqueue with a priority
 - High priority: system processes - sshd, postfix, ...
 - Normal priority: interactive processes - firefox, chrome, ...
 - Low priority: background processes - gcc, make, wget, ...
- Finish all scheduled processes in high priority subqueue, then move to subqueue with higher priority
- Problem:
 - « Starvation »: low priority processes never get CPU

Algorithms

- First Come, First Served
- Shortest-Job-First
- Shortest-Remaining-Time-First
- Round Robin
- Multilevel Queue
- **Multilevel Feedback Queue**

Multilevel Feedback Queue: What?

- An improvement of Multilevel Queue
- Processes can be migrated from a subqueue to another
 - MLQ: processes are **permanently** assigned to subqueue
- Parameters
 - Number of subqueues
 - Scheduling algorithm for each queue
 - Rule to migrate processes

Multilevel Feedback Queue: Example

- « Completely Fair Scheduler »
- « Red-black tree »
- Linux default scheduler, since 2.6.23 (Oct 2007)
- **All** devices
- **All** platforms

Scheduling Algorithms Summary

Algorithm	Preempt?	Priority?	Note
First Come, First Served	No	No	Depends on arrival time
Shortest-Job-First	No	Yes	Low waiting time ω
Shortest-Remaining-Time-First	Yes	Yes	Preemptive SJF, low ω
Round Robin	Yes	No	Low response time ρ
Multilevel Queue	Depends	Depends	Several subqueues, permanent
Multilevel Feedback Queue	Depends	Depends	Several subqueues, migrate

Exercise 2: Scheduling

- Answer the exercise (next slide)
- Write a short report (text file, \LaTeX or Markdown is preferred, **NOT** Word document)
 - Name it « 02.report.scheduling.txt » (or .tex)
 - Write your answers
- Push your report to corresponding forked Github repository

Exercise 2: Scheduling

Process	Arrival Time (ms)	Burst Time (ms)
P_1	0.0	5
P_2	1.0	3
P_3	5.5	2
P_4	6.8	1

Exercise 2: Scheduling

1. Draw Gantt chart of these processes with FCFS, SJF, SRTF, RR ($s = 1ms$)
 - ASCII art is preferred!
2. What is the average waiting time these processes, using...
 - FCFS
 - SJF
 - SRTF
 - RR
3. What is the average turnaround time these processes, using...
 - FCFS
 - SJF

fork-exec

Process Creation

- Start a new process == Create a new process
 - Create new child process
 - Can create child process → grand child process
 - Dependent on OS, parent and child can share
 - **All** resources: opened files, devices, etc...
 - **Some** resources: opened files only
 - **No** resource
- A fully loaded system will have a process tree

Process Creation


```
$ pstree -A
init--+-acpid
      |-cron
      |-daemon---mpt-statusd---sleep
      |-dbus-daemon
      |-dovecot--+-anvil
      |           |-config
      |           `--log
      |-master--+-pickup
      |           |-qmgr
      |           `--tlsmgr
      |-mysqld_safe---mysqld---23*[{mysqld}]
      |-php5-fpm---2*[php5-fpm]
      |-proftpd
      |-screen---bash---python2---{python2}
      |-sshd--+-sshd---sshd---bash---pstree
      |         `--sshd---sshd
      |-udev---2*[udev]
      `--znc---{znc}
```

Process Creation on UNIX/Linux

- New processes are not created from scratch
- Two steps
 - `fork()`
 - `exec()`

Process Creation on UNIX/Linux

- New processes are not created from scratch
- Two steps
 - `fork()`
 - `exec()`



Process Creation on UNIX/Linux

- `fork()`
 - Perfectly «clone» current process to a new process
 - Open files
 - Register states
 - Memory allocations
 - **Except** process id
 - Who's who? Parent? Child?
 - Use return value of `fork()`

```
pid_t fork(void);
```

Process Creation on UNIX/Linux

- Parent: `fork()` returns process id of child
- Child: `fork()` returns 0
- Example

```
#include <unistd.h>
#include <stdio.h>
int main() {
    printf("Main before fork()\n");
    int pid = fork();
    if (pid == 0) printf("I am child after fork()\n");
    else printf("I am parent after fork(), child is %d\n", pid);
    return 0;
}
```

```
$ ./dofork
Main before fork()
I am parent after fork(), child is 2378
I am child after fork()
```

Process Creation on UNIX/Linux

- `exec()`
 - Load an executable binary to replace current process image
 - A family of functions.
 - Ask man

```
int execl(...);
int execlp(...);
int execlp(...);
int execv(...);
int execvp(const char *file, char *const argv[]);
int execvp(...);
```

Process Creation on UNIX/Linux

- exec() example

```
#include <stdio.h>
#include <unistd.h>
int main() {
    printf("Going to launch ps -ef\n");
    char *args[] = { "/bin/ps", "-ef", NULL};
    execvp("/bin/ps", args);
    return 0;
}
```

```
$ ./doexec
```

```
Going to launch ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	0ct19	?	00:01:34	init [2]
root	2	0	0	0ct19	?	00:00:01	[kthreadd]
...							
dovecot	2933	2899	0	0ct19	?	00:00:02	dovecot/anvil
root	2934	2899	0	0ct19	?	00:00:24	dovecot/log
root	3095	1	0	0ct19	?	00:04:57	/usr/lib/postfix/master

Process Creation on UNIX/Linux

```
#include <stdio.h>
#include <unistd.h>
int main() {
    printf("Main before fork()\n");
    int pid = fork();
    if (pid == 0) {
        printf("I am child after fork(), launching ps -ef\n");
        char *args[] = { "/bin/ps", "-ef", NULL };
        execvp("/bin/ps", args);
        printf("Finished launching ps -ef\n");
    }
    else printf("I am parent after fork(), child is %d\n", pid);
    return 0;
}
```

```
$ ./forkexec
```

```
Main before fork()
```

```
I am parent after fork(), child is 12278
```

```
I am child after fork(), launching ps -ef
```

```
UID          PID    PPID    C   STIME TTY          TIME CMD
```


Process Creation on UNIX/Linux

- Questions
 - Why was “Finished launching ps -ef” ’ **not** printed?
 - Why not a single call like Windows `CreateProcess()` or `WinExec()`?
 - Or, why does UNIX separate `fork()` then `exec()`?

Process Creation on UNIX/Linux

- Questions
 - Why was “Finished launching ps -ef” ’ **not** printed?
 - Why not a single call like Windows `CreateProcess()` or `WinExec()`?
 - Or, why does UNIX separate `fork()` then `exec()`?
- Answers
 - Process memory was replaced, the command to print is not there anymore

Process Creation on UNIX/Linux

- Questions
 - Why was “Finished launching ps -ef” ’ **not** printed?
 - Why not a single call like Windows `CreateProcess()` or `WinExec()`?
 - Or, why does UNIX separate `fork()` then `exec()`?
- Answers
 - Process memory was replaced, the command to print is not there anymore
 - `fork()` has different purposes as well, not only in process creation
 - Worker processes (apache, chrome...)

Process Creation on UNIX/Linux

- Synchronization between parent and child
 - `wait()`: wait for termination of a child
 - `waitpid()`: wait for termination of a specific child

Process Termination

- `exit(int status)`
- Cleans up the process and returns to the kernel
- The status is passed to the parent
- Some operating systems also terminate child processes when parent terminates
 - « Cascading termination »

Practical Work 3: mini shell

- Write a new program in C
 - Name it « 03.practical.work.shell.c »
 - Execute commands entered from keyboard
 - Exit shell if user types quit
- Push your C program to corresponding forked Github repository