





# Memory Management

“The memory management on the PowerPC can be used to  
frighten small children.’’

– Linus Torvalds

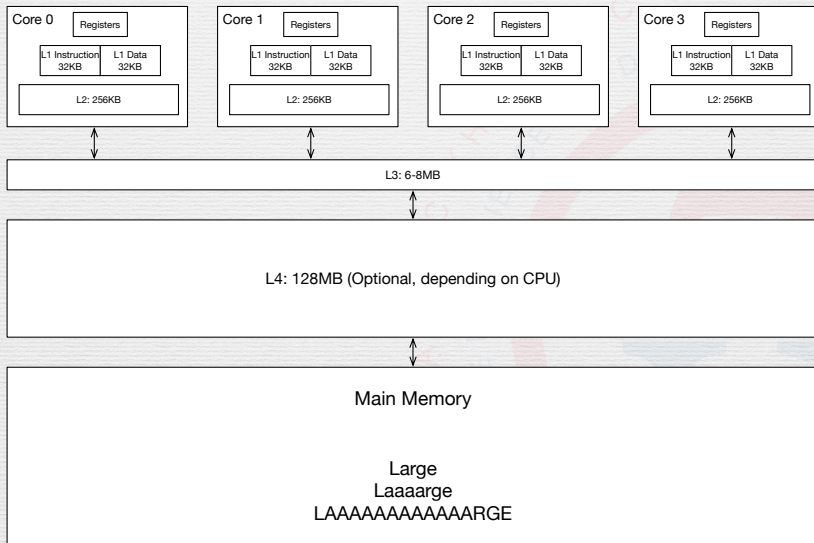
“640 K ought to be enough for anybody.”

- Rumor attributed to Bill Gates, 1981



# Physical Memory

# Registers, Cache and Memory











# Physical Memory

Table 1: Characteristics of different DDR memory<sup>1</sup>

Names	Memory clock	I/O bus clock	Transfer rate	Max bandwidth
SDR-100 <sup>2</sup>	100 MHz	100 MHz	0.1 GT/s	0.8 GB/s
DDR-200	100 MHz	100 MHz	0.2 GT/s	1.6 GB/s
DDR2-800	200 MHz	400 MHz	0.8 GT/s	6.4 GB/s
DDR3-1600	200 MHz	800 MHz	1.6 GT/s	12.8 GB/s
DDR4-3200	400 MHz	1600 MHz	3.2 GT/s	25.6 GB/s

<sup>1</sup>Source: wikipedia and Transend-Info

<sup>2</sup>To be used with Intel Pentium

# Physical Memory

- Similar to a **HUGE** array, shared by everything:
  - Kernel
  - All loaded kernel modules
  - All drivers
  - **All processes**

→ Memory management goal: how to effectively manage physical memory?



# Addressing













The diagram illustrates the mapping between Physical Addressing Space and Logical Addressing Space. It consists of three vertical bars representing memory spaces.

**Physical Addressing Space:** A vertical bar on the left with address markers 0 at the bottom, N in the middle, and max at the top. It is divided into three sections: OS (bottom, 0 to N), Process (middle, N to max), and free (top, max to max).

**Logical Addressing Space:** A vertical bar on the right with address markers 0 at the bottom and max at the top. It is divided into five sections: text (bottom, 0 to max), data (max to max), heap (max to max), free memory (max to max), and stack (top, max to max).

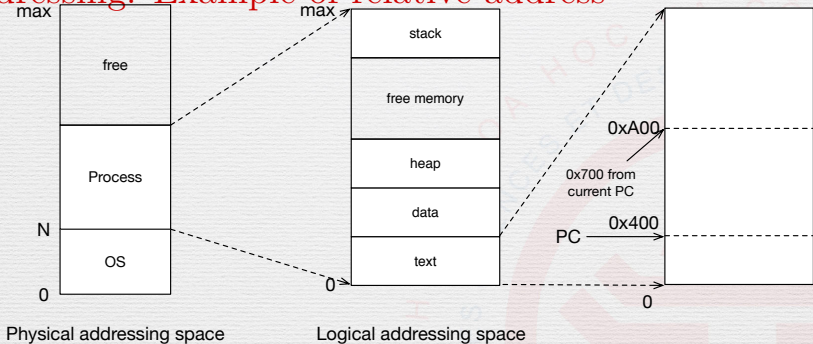
**Mapping:** Dashed arrows show the mapping from the Physical Addressing Space to the Logical Addressing Space. The OS section (0 to N) maps to the text section (0 to max). The Process section (N to max) maps to the data, heap, and free memory sections (max to max). The free section (max to max) maps to the stack section (max to max).

**PC (Program Counter):** A horizontal arrow points to the data section of the Logical Addressing Space, labeled with the address 0x400. Another horizontal arrow points to the top of the data section, labeled with the address 0xA00. A vertical dashed line extends from 0xA00 to the top of the Logical Addressing Space.

**PC (Program Counter):** A horizontal arrow points to the data section of the Logical Addressing Space, labeled with the address 0x400. Another horizontal arrow points to the top of the data section, labeled with the address 0xA00. A vertical dashed line extends from 0xA00 to the top of the Logical Addressing Space.

- 0x700 from current PC in the logical address space
- Question: What is the absolute address of the above address in the physical addressing space?

## Addressing: Example of relative address




- 0x700 from current PC in the logical address space
- Question: What is the absolute address of the above address in the physical addressing space?
  - $N + 0x400 + 0x700$

# Direct Mapping

## Direct Mapping: What?

- Multiprogramming
  - One OS
  - Many processes
- Naive solution: put process memory next to each other in physical memory




## Direct Mapping: Why?

- Supports multiple processes
- The easiest way to share physical memory among OS and processes



## Direct Mapping: Problems

- Protection (Isolation)
- Fragmentation
- Dynamic allocation



## Physical addressing space




## Direct Mapping: Protection

## What?

- OS **protects** its memory and processes memory from illegal access

## Why?

- OS memory should not be accessed by processes
  - System-wide crashes...
- Process A's memory should not be accessed by other processes
  - Privacy and security



## Physical addressing space


## Direct Mapping: Protection

## How to solve?

- Use a **limit** register ( $l$ ) to specify memory range for each process
- Check all memory accesses of process  $i$  against  $l_i$ 
  - Outside range: crash process
- OS is unrestricted: Why?

E.g.

- $l_1 = M - N$
- $l_2 = P - M$
- $l_3 = Q - P$

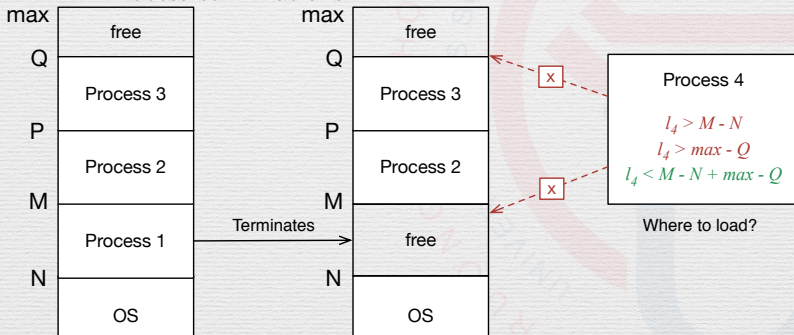


## Physical addressing space

## Direct Mapping: Fragmentation

# What?

- Many small pieces of free memory
  - Process creations
  - Process terminations



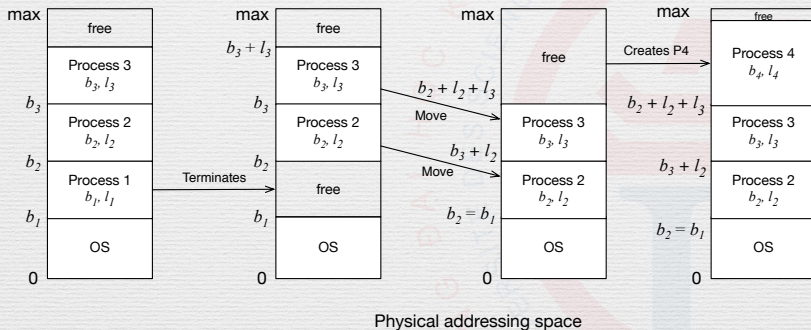
## Direct Mapping: Fragmentation

## How to solve?

- Use a **base** register  $b$  for each process
  - Starting position in memory (base)
  - Use relative addresses like described in “Addressing’’, using  $b_i$
- «Compact» memory regions
  - Make a large free memory block
- Create a new process in this free block


## Direct Mapping: Fragmentation

## How?



## Direct Mapping: Checkpoint

- What?
  - Put process memory next to each other in physical memory
- Why?
  - Naive






## Direct Mapping: Checkpoint

### Problems:

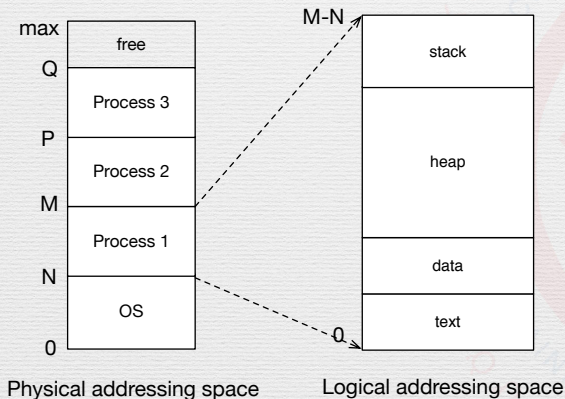
- **Protection (Isolation)**
  - Protects illegal memory access
  - Use **limit** register  $l$
- **Fragmentation**
  - Free memory becomes little pieces
    - Do not fit new processes → Waste
    - Also called «holes»
  - Add **base** register  $b$  with relative address
  - «Compact» memory regions
    - Make a large free memory block
- Dynamic allocation
  - To be continued...







## Direct Mapping: Dynamic Alloc



## Direct Mapping: Dynamic Alloc

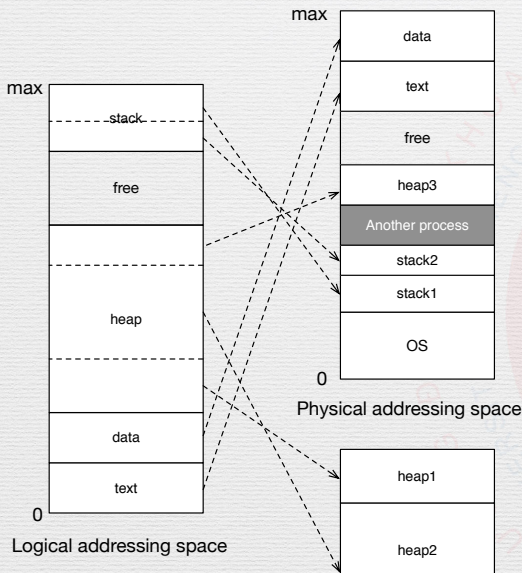
## How to solve?

- Solution 1: Limit amount of memory for each process
  - Consequence 1: fixed heap
    - How will games work?
  - Consequence 2: limit number of concurrent processes at any time
- Solution 2: virtual memory

# Virtual Memory



# Virtual Memory: What?









# Virtual Memory: Why?

- Dynamic allocation
  - More processes concurrency
  - No limit on amount of memory allocation
- Shared memory between processes
- Efficient process creation

# Virtual Memory: How?

- Paging
- Swapping
- Segmentation



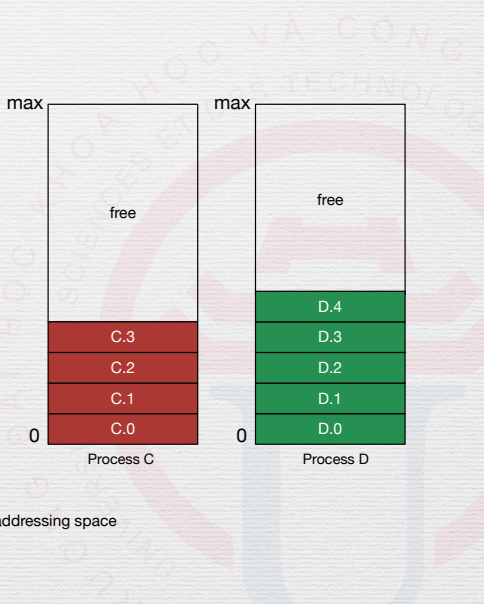


The diagram illustrates the memory layout for two processes, Process C and Process D, within a shared address space. The vertical axis represents the address space, ranging from 0 at the bottom to max at the top.

**Process C:** The memory layout shows a stack of five red blocks labeled C.0, C.1, C.2, and C.3, with an additional unlabeled red block at the bottom. Above these blocks is a large white area labeled "free".

**Process D:** The memory layout shows a stack of five green blocks labeled D.0, D.1, D.2, D.3, and D.4. Above these blocks is a large white area labeled "free".

The diagram demonstrates that the memory layout for Process D is shifted higher in the address space compared to Process C, despite both processes having similar code segments.



# Paging: Why?

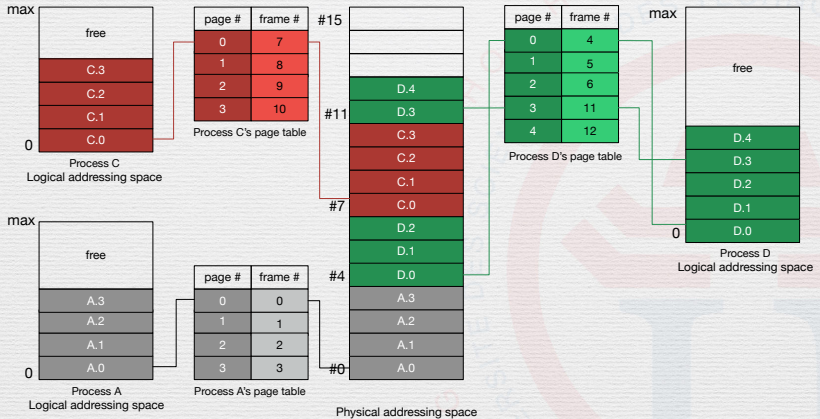
- Allow dynamic allocation
- Allow process memory to be mapped to non-continuous physical memory
- Easy «compactation»








## Paging: How - Page Table



Only to map page, not exact address


## Paging: How - Address Translation

- Split virtual address to 2 parts
  - page number: page id in the page table
  - page offset: distance (how far) from the beginning of this page ( $0 \rightarrow \text{pageSize} - 1$ )
    - i.e. **where** in this page
- Map exact virtual address to exact physical address



## Logical addressing space

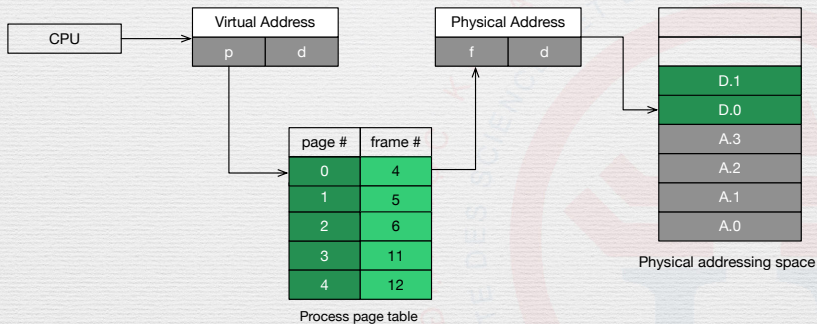
## Paging: How - Address Translation



## Logical addressing space

Physical address  $\rho = \text{pageTable}[p] \times 2^n + d$

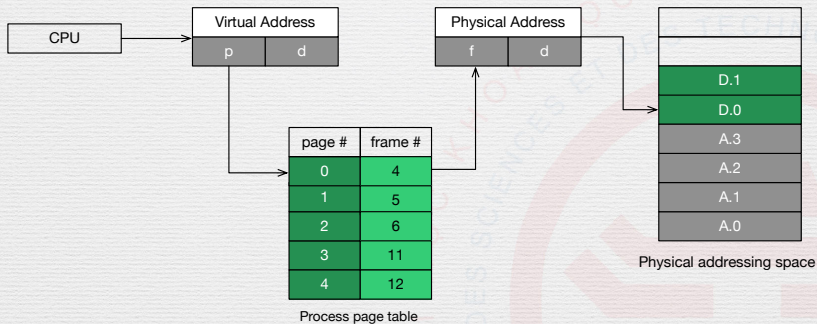
## Paging: How - Address Translation



Physical address  $\rho = \text{pageTable}[p] \times 2^n + d$



## Paging: How - Address Translation Example

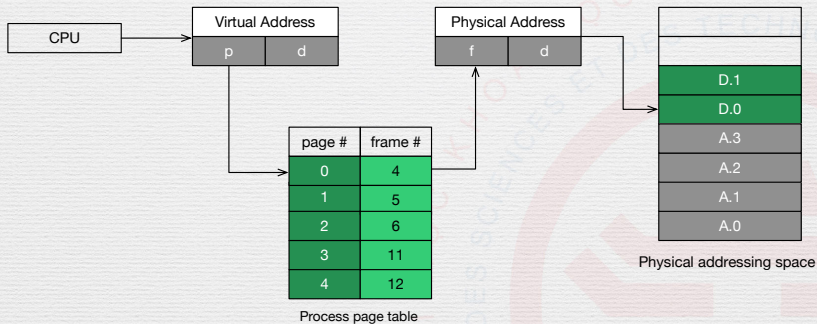


Virtual address  $v = 0x24_{(16)} = 011000_{(2)}$

Virtual address size = 64 ( $= 2^6 \rightarrow m = 6$ )

Page size = 16 ( $= 2^4 \rightarrow n = 4$ )

# Paging: How - Address Translation Example



$$v = 011000_{(2)}, m = 6, n = 4$$

- $\rightarrow d$  has 4 bits
- $\rightarrow p$  has  $6 - 4 = 2$  bits

## Paging: How - Address Translation Example

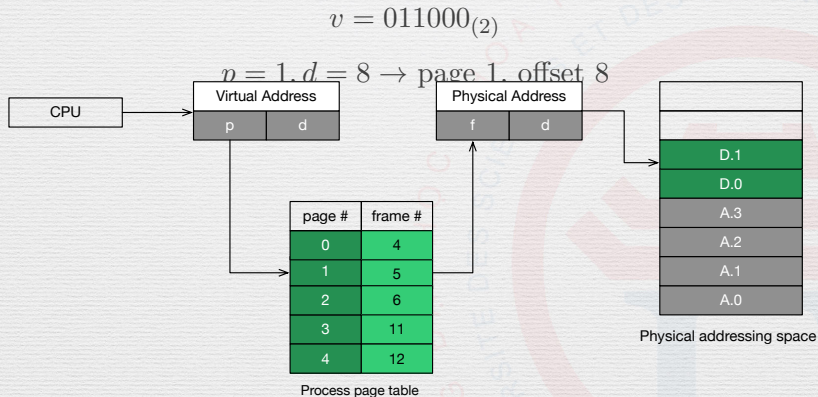
$$v = 011000_{(2)}$$

page number	page offset
01 (2 bits)	1000 (4 bits)

6 bits

$$p = 01_{(2)} = 1_{(10)}, d = 1000_{(2)} = 8_{(10)}$$

# Paging: How - Address Translation Example



$$\text{Physical address } \rho = \text{pageTable}[1] \times 2^4 + 8 = 5 \times 16 + 8 = 88$$

# Paging: Page Table in Reality

- No column “page #”
- One column only
  - Frame id
- Save memory

frame #

# Paging: Advantages

- Solves fragmentation
  - Any free frame can be allocated to a process
- Solves dynamic allocation
  - Process can request as many pages as it needs



# Paging: Problems

- Large
- Performance Inefficient
- Virtual memory must be smaller than physical memory

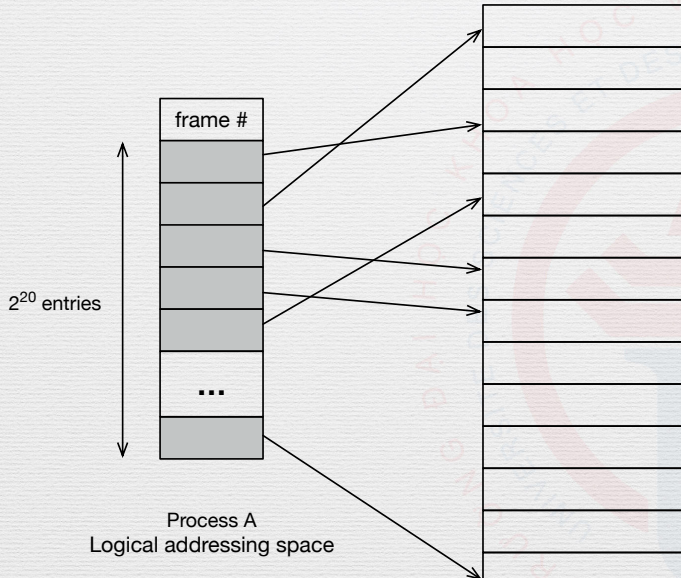
# Paging: Problem 1 - Large

- E.g. 32bit OS: 4GB for virtual memory space
- 4KB page size ( $n = 12$ , or  $d$  is 12-bit)
- Each page table entry is 32-bit
- Question: How big is the page table for **each** process?

# Paging: Problem 1 - Large

- E.g. 32bit OS: 4GB for virtual memory space
- 4KB page size ( $n = 12$ , or  $d$  is 12-bit)
- Each page table entry is 32-bit
- Question: How big is the page table for **each** process?
  - Number of entries:  $4\text{GB} / 4\text{KB} = \frac{4 \times 2^{30}}{4 \times 2^{10}} = 2^{20} = 1048576$
  - Size = number of entries  $\times$  size per entry  
 $= 1048576 \times 4 = 4194304 \text{ bytes} = 4\text{MB}$

# Paging: Problem 1 - Large



# Paging: Problem 1 - Large

On my own MacBook Pro and Hackintosh systems (which are used for writing this slide)

```
$ ps aux | wc -l
```

```
320
```

```
$ echo "`ps aux | wc -l` * 4194304" | bc
```

```
1337982976
```

# Paging: Problem 1 - Large

- Exercise: how large is page table for 64bit OS?
  - Addressing only 48 bits (256TB addressable)
  - 4KB page size
  - Each page table entry is 32 bit
  - Question: How big is the page table for **each** process?



# Paging: Problem 1 - Large

Solution: multi-level page tables


- Split page number to several parts
  - Page directory
  - Page number
- Similar to tree structure

# Paging: Problem 1 - Large

Solution: multi-level page tables

- Split page number to several parts
  - Page directory
  - Page number
- Similar to tree structure
- Page table is paged ☺

# Paging: Problem 1 - Large

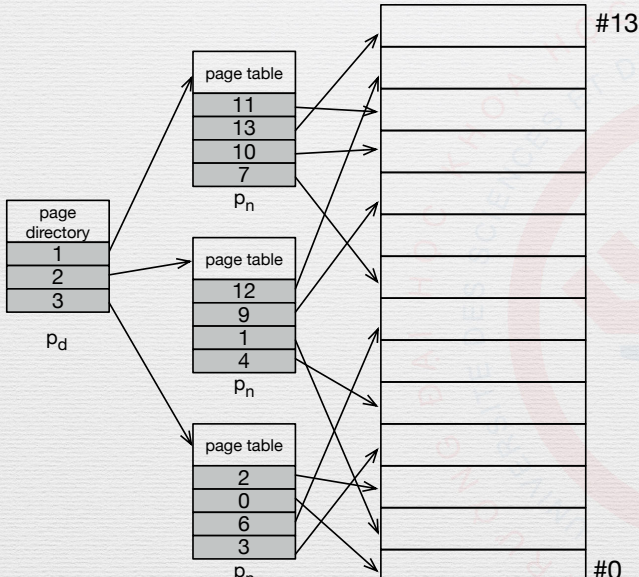


0

# Paging: Problem 1 - Large

- Example for 32-bit OS
  - 32-bit logical address
  - Single level
    - 20-bit page number
    - 12-bit page offset (page size =  $2^{12} = 4\text{KB}$ )
  - Two levels
    - 10-bit page directory
    - 10-bit page number
    - 12-bit page offset

## Paging: Problem 1 - Large

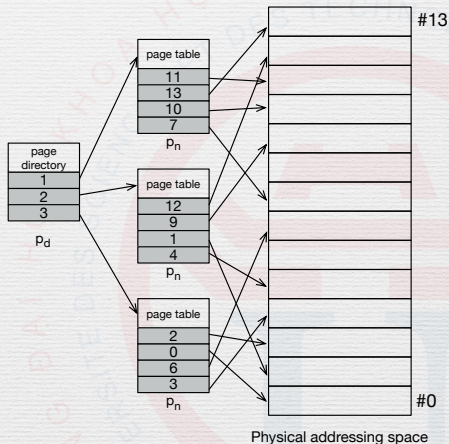


# Paging: Problem 1 - Large

Exercise:

- 2-bit page directory
- 4-bit page number
- 12-bit page offset

What is physical address of  
 $v = 100001110001011100_{(2)}$





## Paging: Problem 1 - Large

- Linux
  - 2.6.9 and below: 3 levels
    - Page Global Directory
    - Page Middle Directory
    - Page Table Entries
    - 512GB addressable
  - 2.6.10 (Jan 2005): 4 levels
    - Page Global Directory
    - Page Upper Directory
    - Page Middle Directory
    - Page Table Entries



# Paging: Problem 1 - Large

- Why is multi-level page table better than single level (flat) one?



## Paging: Problem 2 - Performance Inefficient

- One single memory location access requires 2 memory accesses
  - 1 for accessing page table
  - 1 for the data itself

## Paging: Problem 2 - Performance Inefficient

- One single memory location access requires 2 memory accesses
  - 1 for accessing page table
  - 1 for the data itself
- Question: in Linux x64, how many memory accesses for this C instruction?

```
i++;
```



## Paging: Problem 2 - Performance Inefficient

```
i++;
```

- Answer: 10 physical memory accesses



## Paging: Problem 2 - Performance Inefficient

```
i++;
```

- Answer: 10 physical memory accesses
  - Reading content of `i` to register
    - 1 / 1 / 1 / 1 / 1 for accessing PGD / PUD / PMD / PTE / Content
  - Writing content to `i` after increasing register by 1



## Paging: Problem 2 - Performance Inefficient

Solution:

- «Cache-like» memory for page tables
- Translation Lookaside Buffer (TLB)
- On-chip, single level (flat)
  - Partially parallel
  - Fully parallel
  - Fast!
- Two columns
  - Page id
  - Frame id
- Can be multilevel as well!

## Paging: Problem 2 - Performance Inefficient

Fast, but small <sup>4</sup>

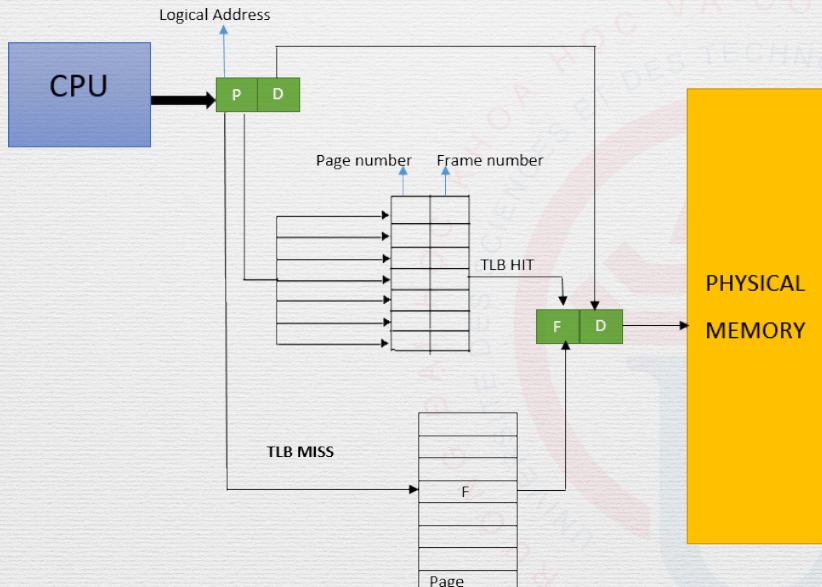
Code name	Generation	TLB L1 entries	TLB L2 entries
Yorkfield	Core 2 Quad	16	256
Nehalem	1 <sup>st</sup> gen, e.g. i7 920	64	512
Ivy Bridge	3 <sup>rd</sup> gen, e.g. i7 3632QM	64	512
Haswell	4 <sup>th</sup> gen, e.g. i7 4770HQ	64	1024
Skylake	6 <sup>th</sup> gen, e.g. i7 6700K	64	1536
Kaby Lake	7 <sup>th</sup> gen, e.g. i7 7700K	64	1536

<sup>4</sup>Data courtesy of 7-CPU





## Paging: Problem 2 - Performance Inefficient



## Paging: Problem 3 - Small Virtual Memory

Remind:

- Virtual Memory > Physical Memory
- Code / data must be in physical memory

→ temporarily use disk as storage for unused/least used pages





## Paging: Problem 3 - Small Virtual Memory

- When to swap to disk?
  - Low on memory
  - Very few access on a page
- When to swap back?
  - Access to a swapped page
  - «Page fault»



# Memory Management

“The memory management on the PowerPC can be used to  
frighten small children.’’

– Linus Torvalds

# Memory Management

“The memory management on the PowerPC can be used to  
frighten small children.’’

– Linus Torvalds

“640 K ought to be enough for anybody.”

- Rumor attributed to Bill Gates, 1981